

The Abstraction: The process

What is a process?

Reading

Operating Systems: Three Easy Pieces. Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau.


Chapter 4: The Abstraction: The Process

Chapter 5: Interlude: Process API

Program vs process


A *program* is instructions stored on disk – an executable file

A *process* is a running program




Circles in a Circle, Wassily Kandinsky, 1923 [\[source\]](#)

Process Creation



Address Space



The *address space* is where the programs data resides.

Process Information

A process is more than a list of instructions, execution requires knowing...

Which instruction to execute next? *program counter*

What is the immediate data on which instructions operate? *registers*

Where are the instructions and data stored in memory? *address space*

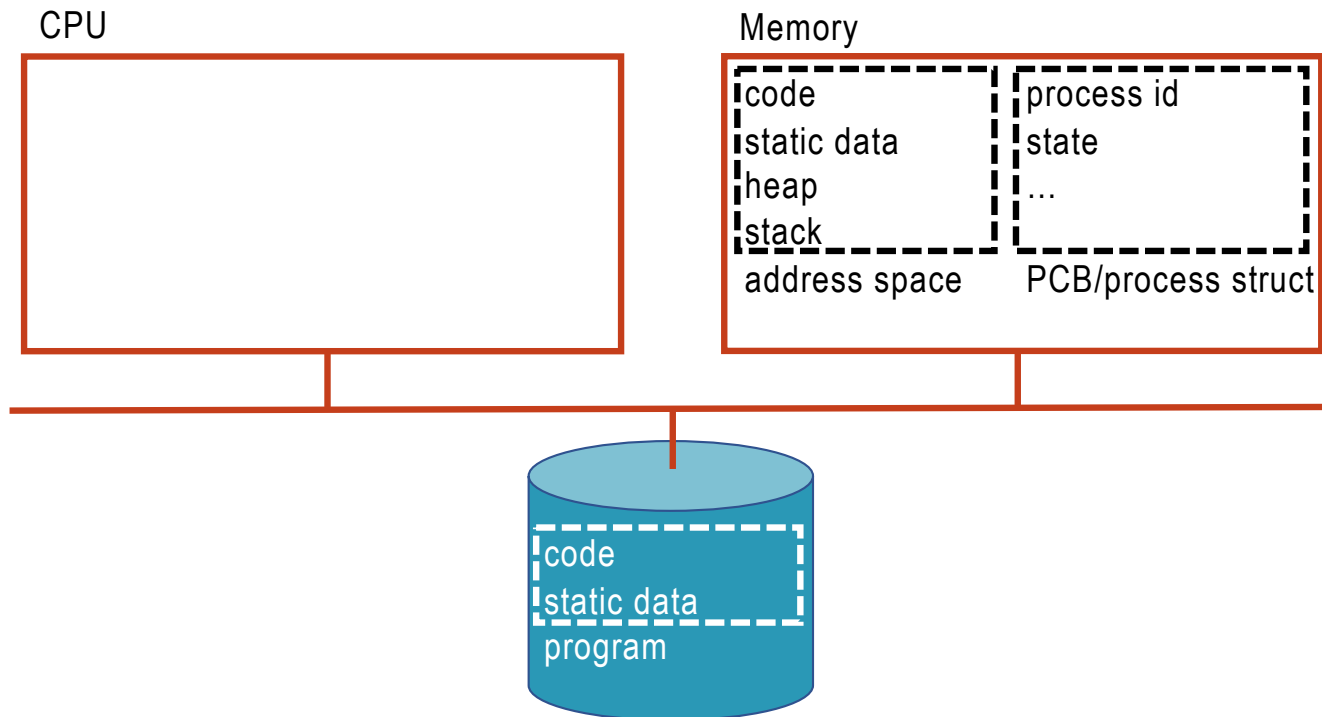
How to find parameters of current function? *stack and frame pointer*

And more that we will cover later in this class (e.g., open files, threads...)

Most of these terms should be familiar, if not,


time to review your COM S 321/CPR E 228 notes.

Process Control Block



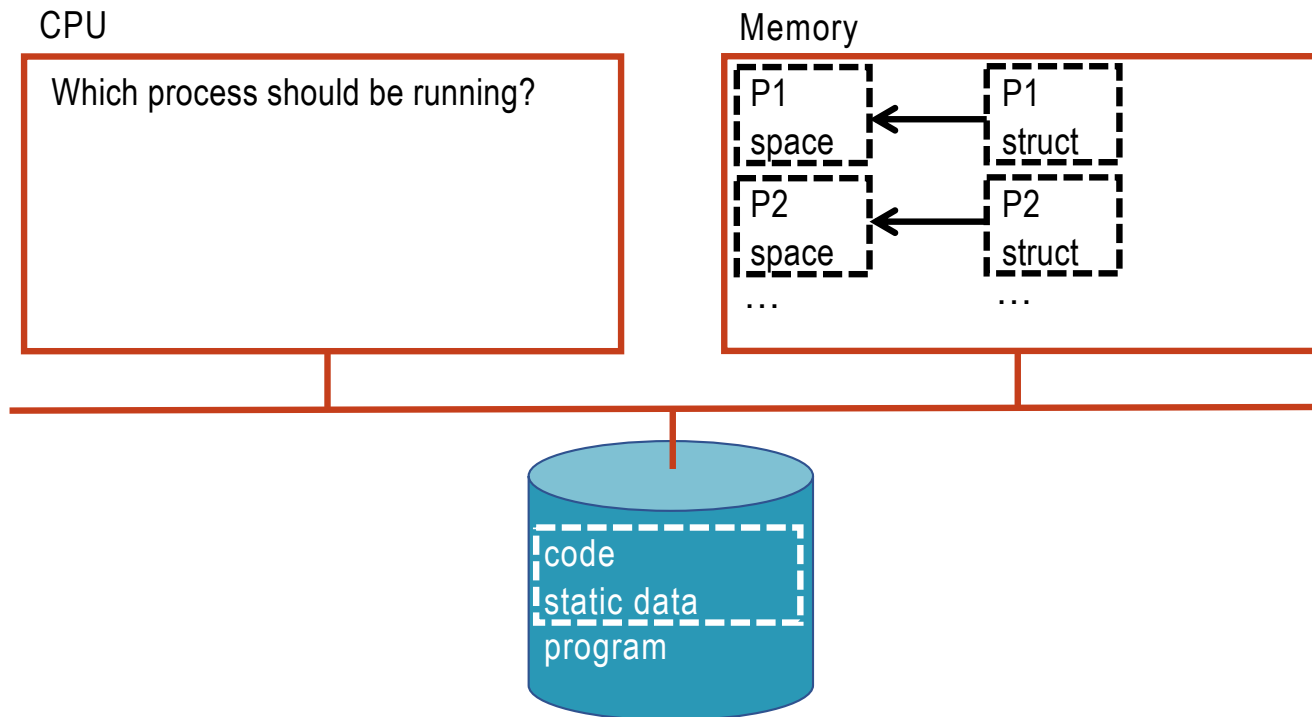
The *Process Control Block (PCB)* is a structure the OS uses to keep track of the process information.

Process In Execution




During execution, the state of the CPU is dedicated to the process.

Multiple Processes



We don't want to be limited to one process, but a CPU can only execute one process at a time...

Scheduler: Running and Ready



The **scheduler** is part of the OS, its job is to manage the process' **states**
Only one process **running** at a time (single core assumption)
Other processes that are ready to execute go in a **ready queue**

Multiprogramming

Processes need to perform I/O, e.g., disk read/write

Consider a process the makes a random read from disk

Hard drive latency 10ms and SSD latency 1ms

CPU clock 0.3ns cycle and 1 instructions per cycle

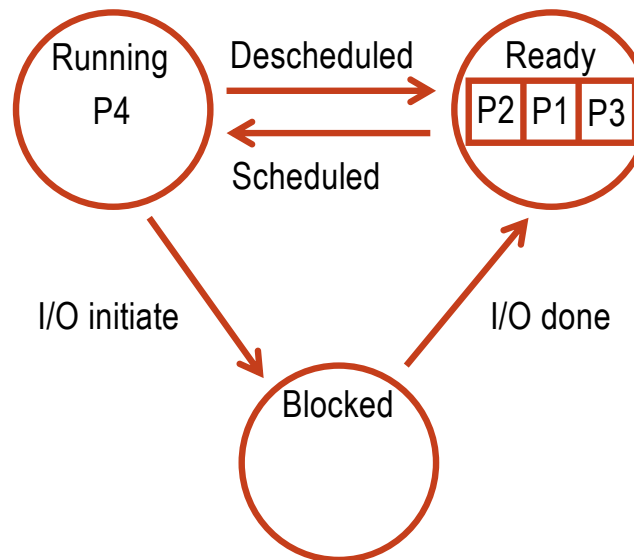
10 million instructions wasted waiting for SSD and 100 million for hard drive!!!!

Clearly not good use of the CPU

Multiprogramming is letting the user run multiple programs (processes) together

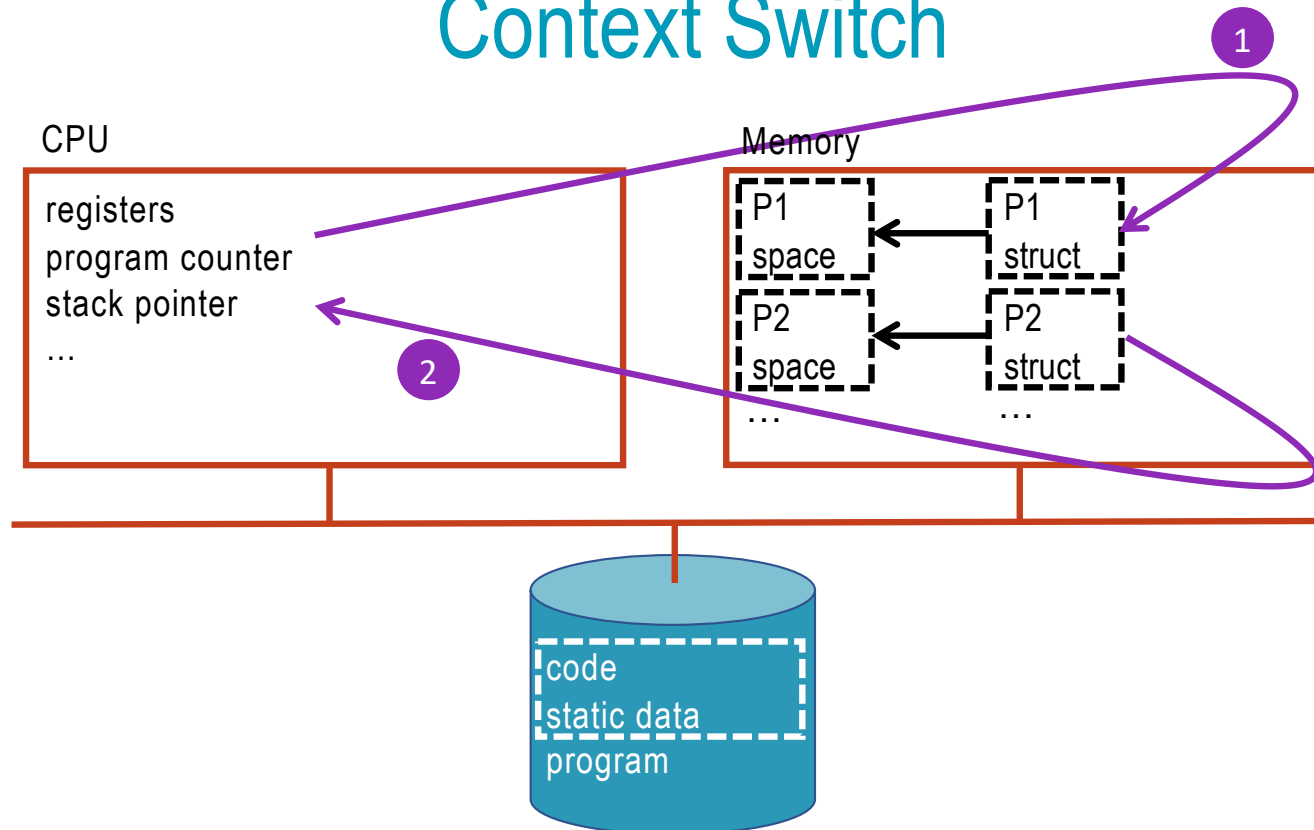
When one process needs to wait for I/O another can be scheduled on the CPU

Scheduler with Blocked State



When a process is waiting for I/O it is placed in the *blocked* state
Next process in ready queue is set to running

Context Switch



Save CPU context of descheduled process
Load CPU context of scheduled process

Process Tree

Every process is created by another process (except to root process)

Every process must have a *parent*, by default the parent is the process that created it (the *child* process)

When a process is created it starts in an *initial state* and when it terminates it goes to a *final (zombie) state*

The final state is required because we can't delete the PCB immediately, the parent may want information such as the exit code

If a parent terminates before the child, the child becomes an *orphan* process and the root process becomes its parent

Process Control Block (struct proc) in xv6

proc.h

```
// Saved registers for kernel context switches. // Per-process state
struct context {
    uint64 ra;
    uint64 sp;

    // callee-saved
    uint64 s0;
    uint64 s1;
    uint64 s2;
    uint64 s3;
    uint64 s4;
    uint64 s5;
    uint64 s6;
    uint64 s7;
    uint64 s8;
    uint64 s9;
    uint64 s10;
    uint64 s11;
};

enum procstate { UNUSED, USED, SLEEPING,
RUNNABLE, RUNNING, ZOMBIE };

struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum procstate state; // Process state
    void *chan; // If non-zero, sleeping on chan
    int killed; // If non-zero, have been killed
    int xstate; // Exit status to be returned to parent's wait
    int pid; // Process ID

    // proc_tree_lock must be held when using this:
    struct proc *parent; // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack; // Virtual address of kernel stack
    uint64 sz; // Size of process memory (bytes)
    pagetable_t pagetable; // User page table
    struct trapframe *trapframe; // data page for trampoline.S
    struct context context; // swtch() here to run process
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
};
```

Process System Calls

POSIX (Portable Operating System Interface)

Standard programming interface provided by UNIX like systems

fork()

Create another process (child) that is a copy of the current process (parent)

exec()

Change the program of the currently executing process

wait()

Do nothing until a child process has terminated



Tuning fork [\[source\]](#)

fork() // Process Creation

```
pid_t fork(void);
```

Creates a new process by duplicating the calling process

Child process has a copy of parent's address space

On success:

- Both parent and child continue execution at the point of return from fork()

- Returns pid of the child process to the parent process; returns 0 to child process

On failure:

- Child is not created; returns -1 to parent

fork.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  int main(int args, char *argv[]) {
6      printf("hello world (pid:%d)\n", (int) getpid());
7      int rc = fork();
8      if (rc < 0) { // fork failed; exit
9          fprintf(stderr, "fork failed\n");
10         exit(1);
11     } else if (rc == 0) { // child (new process)
12         printf("hello, I am child (pid:%d)\n", (int) getpid());
13     } else { // parent
14         printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
15     }
16     return 0;
17 }
```

console

```
hello world (pid:19979)
hello, I am parent of 19980 (pid:19979)
hello, I am child (pid:19980)
```

parent

```
7 int rc = fork();  
8 if (rc < 0) { // fork failed; exit  
9     fprintf(stderr, "fork failed\n");  
10    exit(1);  
11 } else if (rc == 0) { // child (new process)  
12     printf("hello, I am child (pid:%d)\n", (int) getpid());  
13 } else { // parent  
14     printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());  
15 }  
16 return 0;  
17 }
```


child

```
8 if (rc < 0) { // fork failed; exit  
9     fprintf(stderr, "fork failed\n");  
10    exit(1);  
11 } else if (rc == 0) { // child (new process)  
12     printf("hello, I am child (pid:%d)\n", (int) getpid());  
13 } else { // parent  
14     printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());  
15 }
```



Process Tree

Time


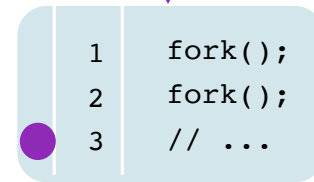


1



2



3




What Happens First?


Time

1

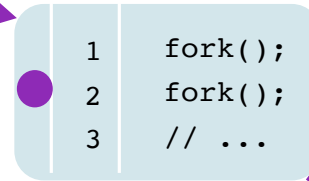
A




2




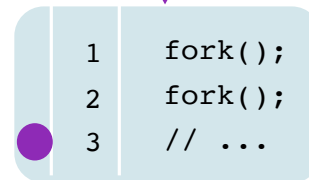

C



3



B



Question: What is the correct order for A, B and C?

wait() // Wait for Child

```
pid_t wait(int *wstatus);
```

Suspend execution of the parent until one of its children terminates

On success:

- Returns pid of the child process that terminated

- wstatus is populated with information about the way the child process terminated

If a process has terminated, but parent has not yet called wait(), the process becomes a *zombie*

If the parent terminated without calling wait(), the child process becomes an *orphan*, and a system process (systemd in Linux) becomes the parent

exec() // Change the Program

```
int exec(const char *pathname, char *const argv[]);
```

Replaces the current program with a new one

Command line arguments are passed in argv

The process keeps its file descriptors

On success:

- The process is running a new program

- The function does not return (no where to return)


On failure:

- The function returns -1

First argument is always the name of the executable

Example usage:

```
char *args[] = { "wc", "README", 0 };  
exec( "wc", args );
```



View Processes on Command Line

```
$ ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1  0.0  0.1 186896 15884 ?        Ss   2020   13:22 /usr/lib/systemd/systemd --system --deserialize 39
root    2758580  0.1  0.1  41896   9916 ?        Sns  00:45    0:00 sshd: tancreti [priv]
tancreti 2758586  0.0  0.0  41772   5904 ?        RN   00:45    0:00 sshd: tancreti@pts/1
root    2758587  0.0  0.0      0      0 ?        I    00:45    0:00 [kworker/3:2]
tancreti 2758588  0.5  0.0  17352   5444 pts/1    Sns  00:45    0:00 -bash
root    2758598  0.0  0.0      0      0 ?        I    00:45    0:00 [kworker/4:2]
tancreti 2758634  0.0  0.0  17932   3616 pts/1    RN+  00:45    0:00 ps aux
```

systemd, Linux parent, PID=1

SSH processes managing my connection

bash shell process I am interacting with

ps command (it saw itself)

Exploring Process Tree on Command Line

```
$ pstree
systemd--NetworkManager--2*[ {NetworkManager} ]
|--sshd--sshd--sshd--bash--pstree
|      |--sshd--sshd--bash
...
```

pstree command

bash shell

SSH child to manage connection A

SSH child to manage connection B

SSH server

systemd is Linux parent process

Question: How did bash create a child process that executes pstree?

Interprocess Communication

A pipe is a unidirectional data channel that can be used to communicate from one process to another

- Sender puts data to one end (the write-end of the pipe)

- Receiver gets data from the other end (the read-end of the pipe)

Common example is the shell, it creates two children and pipes standard out (e.g., `printf`) of one into standard in (e.g., `read`) of other

Pipe Creation




Figure 15.3 Half-duplex pipe after a fork

pipe() // Connect two processes

```
int pipe(int p[2]);
```

Creates communication channel

Typical usage is right before calling fork, each process must close the ends of the pipe it is not using

On success:

p[0] is file descriptor of read side of pipe, p[1] is write side of pipe

Returns 0

On failure:

Returns -1

Example Usage:

See sh.c

dup() // Duplicate file descriptor

```
int dup(int fd);
```

Returns new file descriptor that is the lowest numbered available descriptor

New file descriptor refers to the same source as fd did previously

For example, closing standard out (1) and then calling dup(fd) will cause all calls to printf to be directed to what fd pointed to

On success:

- New file descriptor points to source of provided file descriptor

- Returns new file descriptor

On failure:

- Returns -1

Example Usage:

- See sh.c