

The Abstraction: Address Spaces

How do programs share main memory?

By Matthew Tancreti
For COM S 352
Iowa State University

Address Spaces

Processes must share main memory, we must account for processes

created and destroyed

requesting more memory during execution


requesting more memory than is physically available

How to meet the memory demands
of processes?



Premier Disque, Robert Delaunay, 1913 [\[source\]](#)

Process Address Space



A process' view of memory is called its **address space**

The address space is an abstraction of the physical memory, assumptions:

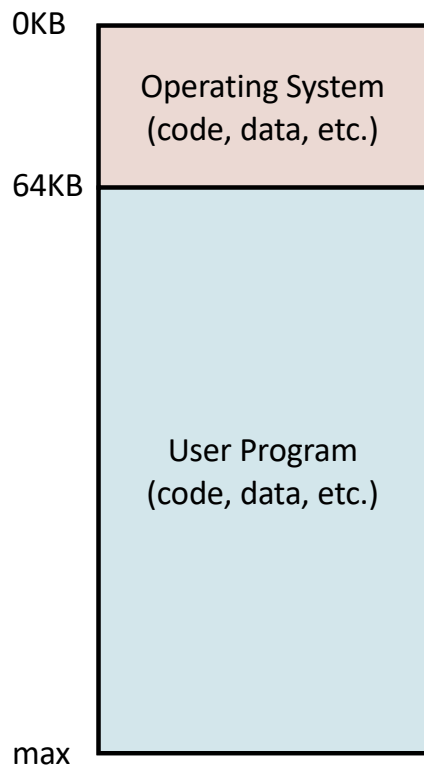
- Address space starts at 0
- Address space is contiguous
- All address available at any time

At the top are static items (e.g., code, global variables and constants)

Processes have **two forms of dynamic memory**: **heap** and the call **stack**

Because they both need to grow an unknown amount, it is logical to place them on opposite ends of the address space

Single Program

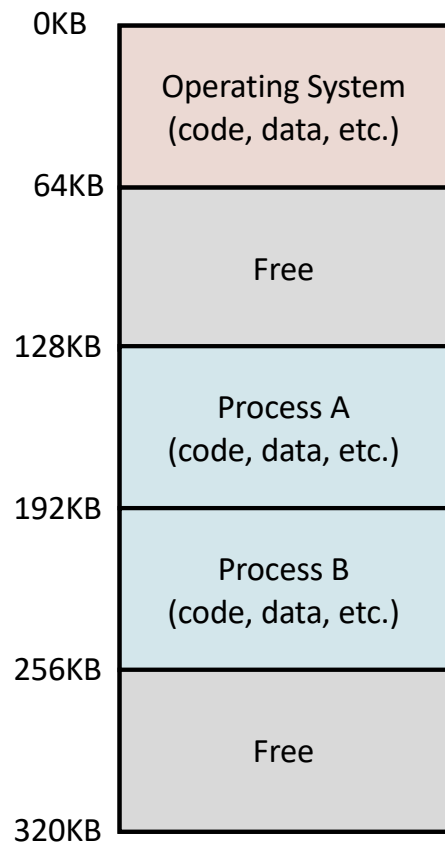


Main memory divided between the OS and user program

Addresses assigned statically (at compile time)

Still used in simple (single program) embedded systems

Multiprogramming

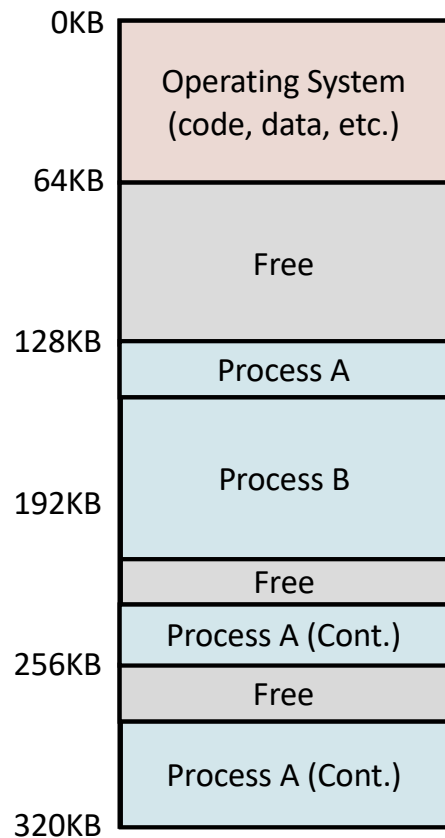


Processes share main memory

Simplest scheme is to assign contiguous regions of memory

Becomes costly when a process needs to grow its memory

Memory Virtualization




We want to give users an easy view of memory, but the reality is much more messy

Memory can become fragmented

Unused address space doesn't need to be mapped to physical memory

Not every process will fit in memory, use disk for extra storage

Called **memory virtualization**



GB vs GiB

Early programmers noticed an approximate relationship between binary and decimal numbers

$2^{10} = 1,024$ which is almost 1 thousand

$2^{20} = 1,048,576$ which is almost 1 million...

They adapted a binary version of the SI units

$2^{10} = 1 \text{ KB}$

$2^{20} = 1 \text{ MB}$...

This became confusing when marketing memory and disk sizes, so new SI binary units were created

In this class we will use the traditional interpretation (e.g., $1\text{KB} = 1,024$, $1\text{MB} = 1,048,576$...)

SI Decimal		SI Binary	
10^3	kilobyte (KB)	$2^{10}=1024$	kibibyte (KiB)
10^6	megabyte (MB)	$2^{20}=1,048,576$	mebibyte (MiB)
10^9	gigabyte (GB)	$2^{30}=1,073,741,824$	gibibyte (GiB)
10^{12}	terabyte (TB)	$2^{40}=1,099,511,627,776$	tebibyte (TiB)

Conversions

Problem: write 62,000,000,000 bytes in SI notation


$62,000,000,000 / 2^{30} = 57.74\text{GiB}$ (or in this class GB, assuming binary interpretation)

Problem: how many bytes is 32KB (assuming binary interpretation of KB)?


$32 * 2^{10} = 32,768$ bytes

Word Size



Processor architecture determines word size



Registers, instructions, address bus and data bus are typically one word




Example: 32-bit machine can address 4GB of memory



Masking

Making is used to reveal any number of specific bits

Question: How to get the value of only to 4 least significant digits?




```
result = data & 0x0F
```

Flags

Flags are a single bit to represent a Boolean value

Used in many encoding schemes (e.g., compression, images, instructions, etc.)

For example: Suppose bit 7 indicates extended mode, how to test bit 7?



```
if (data & 0x80)
```