

Paging

How to make physical memory more flexible?

By Matthew Tancreti
For COM S 352
Iowa State University


Paging

Segmentation made code, stack and heap independently relocatable

Segments can become arbitrarily large contiguous regions of memory

Resulted in **external fragmentation**

What if we divide the address space into equal size pages?



Movable type printing invented by Bi Sheng 1031-1095 [\[source\]](#)

Paging

Address space divided into equal sized pages that can be stored in frames

Address Space	Virtual Addresses
page 0	0
page 1	16K
page 2	32K
page 3	48K
page 4	64K
page 5	80K
page 6	96K
page 7	112K
	128K

Physical Addresses	Physical Memory	
0	Operating System	frame 0
16K	(unused)	frame 1
32K	page 7	frame 2
48K	page 0	frame 3
64K	(unused)	frame 4
80K	page 1	frame 5
96K	(unused)	frame 6
112K	page 2	frame 7
128K		

Page Table

One page table for each process

Virtual Page Number (VPN) is the index of the table


Physical Frame Number (PFN) points to the frame in physical memory

Valid bit indicates if table entry is valid (not all of address space needs to be mapped)

Address Space	Virtual Addresses	Page Table			Physical Addresses	Physical Memory	
	0	VPN	PFN	valid	0		
page 0	16K	0	3	1	16K	Operating System	frame 0
page 1	32K	1	5	1	32K	(unused)	frame 1
page 2	48K	2	7	1	48K	page 7	frame 2
page 3	64K	3	-	0	64K	page 0	frame 3
page 4	80K	4	-	0	80K	(unused)	frame 4
page 5	96K	5	-	0	96K	page 1	frame 5
page 6	112K	6	-	0	112K	(unused)	frame 6
page 7	128K	7	2	1	128K	page 2	frame 7

Virtual Address Bits


Virtual address divided into VPN and offset



```
VPN = VirtualAddress >> NUM_OFFSET_BITS
```

```
Offset = VirtualAddress & OFFSET_MASK
```

Virtual Address VPN Bit Size



If total address space is 2^m bytes and page size is 2^n bytes then

VPN is $m-n$ bits


offset is n bits

Question: If address space is 1GB and page size is 4KB how many bits is the VPN?

Answer: $1\text{GB} = 2^{30}$, $4\text{KB} = 2^{12}$, therefore VPN is $30 - 12 = 18$ bits

Check: $4\text{KB} * 2^{18} = 1\text{GB}$

Physical Address Bits



If physical memory is 2^m bytes and frame (also page) size is 2^n bytes then

PFN is m-n bits

offset is n bits

Therefore, calculate physical address as:

$$\text{PhysAddr} = \text{PFN} * \text{frame_size} + \text{offset}$$

or in binary arithmetic:

$$\text{PhysAddr} = (\text{PFN} \ll \text{NUM_OFFSET_BITS}) \mid \text{offset}$$

Keep in mind, total physical memory and address space size may not be the same


Question: Assume the page size is 4KB, what is the address for frame 10 and offset 128?

Answer: $4\text{KB} * 10 + 128 = 41,088$

Address Translation

Steps performed **in hardware** (MMU):

1. Compute page number
 $\text{VPN} = \text{VirtualAddress} \gg \text{NUM_OFFSET_BITS}$
2. Look up page in page table to find frame number
3. Confirm that access to page is allowed (e.g., valid bit set)
4. Compute physical address
 $\text{PhysAddr} = (\text{PFN} \ll \text{NUM_OFFSET_BITS}) \mid \text{offset}$



Unused Pages


Unused pages don't need to be mapped to physical memory

Process is unaware of pages, when it requests more memory allocated for heap or stack the OS maps pages as needed

Address Space	Virtual Addresses	Page Table			Physical Addresses	Physical Memory	
	0	VPN	PFN	valid	0	Operating System	frame 0
page 0	16K	0	2	1	16K	(unused)	frame 1
page 1	32K	1	5	1	32K	page 7	frame 2
page 2	48K	2	7	1	48K	page 0	frame 3
(unused)	64K	3	-	0	64K	(unused)	frame 4
(unused)	80K	4	-	0	80K	page 1	frame 5
(unused)	96K	5	-	0	96K	(unused)	frame 6
(unused)	112K	6	-	0	112K	page 2	frame 7
page 7	128K	7	2	1	128K		


Free Memory Management

OS needs to know which frames are not in use, simple method is linked list



No external fragmentation - no unusable gaps between frames

Pages can have **internal fragmentation** - unused portion of page



Page Size

Tradeoff

Small page size means bigger page table (more page table entries)

Bigger page size means more internal fragmentation

Typical page size is 8KB in Linux

Concern

Each process has its own page table which the OS stores in frames in physical memory

Linux page size is 8KB, on pyrite we saw the address space is 140TB
~17 million page table entries for every process!

Linux reduces size with multi-level (3-level tree structure) page tables

Page table lookup is slow, every memory access requires additional memory access(es)! How to speed up memory?