

# SSD

How to build a flash-based SSD?

By Matthew Tancreti  
For COM S 352  
Iowa State University


# SSD

**Solid-State Drive (SSD)** is made from **flash memory**, a silicon transistor technology

Unlike HDD there are no moving parts – no spinning platters or mechanical arm

After decades, SSD becoming more popular than HDD, what are the consequences on file systems?

How build a flash-based SSD?



SSD [\[source\]](#)

# NAND-based flash

A transistor forms a **cell** which stores a bit in a single-level cell or up to three bits in a trip-level cell

To write to a value the cell must first be erased, setting all bits to 1

Only after a bit is set to 1 can it be set to 0

The need to **erase before write** is important to understanding SSD performance characteristics

# Data Organization


Hierarchically organized

**Bank** consists of multiple blocks

**Block** is typically 128KB and consists of multiple pages

**Page** is typically 4KB

Individual bits stored in **cells**



Warning: the words page and block have a different meaning for SSD than they have in other contexts, such as virtual memory or HDD

# Basic Flash Operations

Flash operations are performed at the page/block level

**Read (page)** – client provides page number to read, relatively fast (e.g., 25 $\mu$ s), does not depend on location of page or previous page (**random access** device)

**Erase (block)** – before programming the block must be set to all 1, orders of magnitude slower than read (e.g., 1.5ms)

**Program (page)** – writes the page by setting 1's to 0's where needed, time is somewhere between a read and erase (e.g., 200 $\mu$ s)

# Example of Basic Flash Operations

Assume we want to write to page 0.

Starting state

Page 0	Page 1	Page 2	Page 3
00011000	11001110	00000001	00111111
VALID	VALID	VALID	VALID

← Need to figure out what to do with pages 1 to 3, e.g., move them to another block.

Erase block

Page 0	Page 1	Page 2	Page 3
11111111	11111111	11111111	11111111
ERASED	ERASED	ERASED	ERASED

Program page 0

Page 0	Page 1	Page 2	Page 3
00000011	11111111	11111111	11111111
VALID	ERASED	ERASED	ERASED

## HDD vs SSD Latencies

Rotational delay – time for sector to rotate under the disk head, 66 $\mu$ s for one rotation on fastest drives

Seek time – time for disk arm to change position to the correct track, 4ms typical

Read (page) – fast and random access, 25 $\mu$ s

Erase (block) – 1.5ms

Program (page) – 200 $\mu$ s

These are only latencies, must also consider transfer times

# Wear Out

**Wear out** is the issue that a flash block has a limited number of times (100,000) it can be erased before it becomes unusable

Suppose a program writes 100,000 times to a file, seems like a show stopper

The SSD mitigates this problem with **wear leveling**

- Logic page address is independent of physical page on flash

- Every time a page is written a new physical page is used to store the logic page

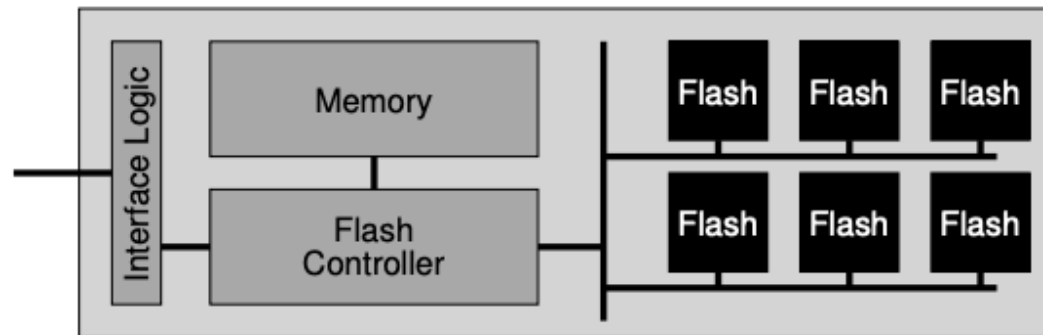
- In this way a logic page can be written to 100,000 times but every write is to a different physical location



# Logical Diagram of SSD

A flash controller virtualizes the flash (i.e., mapping logical address to virtual address) in the **Flash Translation Layer (FTL)**

On chip memory is required for caching and buffering of blocks and pages



# Log-Structured FTL

## Motivations

- Large time cost of erasing a block before pages can be written

- Want to have wear leveling (independence between virtual and physical addresses)


In **log-structured FTL** an in-memory table is used to map virtual to physical pages

On every write the page is moved to a different physical location

# Log-Structured Example

```
Write(100) with contents a1
Write(101) with contents a2
Write(2000) with contents b1
Write(2001) with contents b2
```

Block starts with all invalid

Block:	0				1				2			
Page:	00	01	02	03	04	05	06	07	08	09	10	11
Content:												
State:	i	i	i	i	i	i	i	i	i	i	i	i

## Erase block

Block:	0				1				2			
Page:	00	01	02	03	04	05	06	07	08	09	10	11
Content:												
State:	E	E	E	E	i	i	i	i	i	i	i	i

Write a1 to first free physical  
Page and log mapping

Table:	100 → 0												Memory
Block:	0				1				2				
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a1												
State:	V	E	E	E	i	i	i	i	i	i	i	i	
													Flash Chip

Write other pages to free  
Physical pages and log  
mappings

Table:	100 → 0	101 → 1	2000 → 2	2001 → 3	Memory
Block:					
Page:					
Content:	a1	a2	b1	b2	
State:	V	V	V	V	

# Garbage Collection

If the same logical page is written multiple times, the old versions of the page will remain in physical memory as garbage (unusable)

Table:	100	→	4	101	→	5	2000	→	2	2001	→	3	Memory
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a1	a2	b1	b2	c1	c2							
State:	V	V	V	V	V	V	E	E	i	i	i	i	

**Garbage collection** is the reclamation of dead blocks

In example above, b1 and b2 can be moved to physical pages 6 and 7, then block 0 can be erased for reuse

# Mapping Table Size

Mapping table can be very large

Assume 1TB SSD and 4 byte entry for each 4KB page, then map is 1GB

A **hybrid mapping** approach can map at either the page or block level, far fewer blocks so less mapping required

## Example Performance of HDD vs SSD

Device	Random		Sequential	
	Reads (MB/s)	Writes (MB/s)	Reads (MB/s)	Writes (MB/s)
Samsung 840 Pro SSD	103	287	421	384
Seagate 600 SSD	84	252	424	374
Intel SSD 335 SSD	39	222	344	354
Seagate Savvio 15K.3 HDD	2	2	223	223