

1. (20 points) Consider a database that has two integer objects, A and B. DBMS is given two programs:

P1	P2
$A = A + 1$ $B = B + 1$	$A = B$ $B = B + 1$

Before the program is execution, $A = 0$ and $B = 0$. For each of the following schedules, determine if it is a serializable schedule. You must explain why.

S1

T1	T2
$A = A + 1$ $B = B + 1$	$A = B$ $B = B + 1$

For this problem, there should be two serial schedules.

1. Execute P1 and then P2. The result is $A = 1$ and $B = 2$
2. Execute P2 and then P2. The result is $A = 1$ and $B = 2$

S1 is a serializable schedule. The result of S1 is $A = 1$ and $B = 2$ which is equivalent to one of the two serial schedules.

S2

T1	T2
$A = A + 1$ $B = B + 1$	$A = B$ $B = B + 1$

S2 is not a serializable schedule. The result of S2 is $A = 0$ and $B = 2$ which is not equivalent to one of the two serial schedules.

S3

T1	T2
$A = A + 1$ $B = B + 1$	$A = B$ $B = B + 1$

S3 is a serializable schedule. The result of S3 is $A = 1$ and $B = 2$ which is equivalent to one of the two serial schedules.

2. (20 points) Consider the following two programs.

P1	P2
R(A)	W(A)
R(C)	R(B)
W(C)	W(B)
c/o	c/o

For each of the following execution, determine if it is a schedule, a serial schedule, and/or a strict schedule. You must explain why.

S1		S2		S3		S4	
T1	T2	T1	T2	T1	T2	T1	T2
R(A)	W(A)	R(A)	W(A)	R(A)	c/o	R(A)	W(A)
R(C)		R(C)		R(C)			
W(C)	R(B)	W(C)		W(C)		W(C)	R(B) W(B) c/o
c/o	W(B)	c/o		c/o		c/o	
	c/o		R(B)		W(A)		
			W(B)		R(B)		
			c/o		W(B)		
					c/o		

S1. The given schedule is a schedule. It reads and writes in the same order as in their original programs. This is not a serial schedule because the execution of T1 and T2 interleaves. This is not a strict schedule because T1 has an R(A), before T1 commits/aborts, W(A) is T2 is executed.

S2. The given schedule is a schedule. It reads and writes in the same order as in their original programs. This is not a serial schedule because the execute of T1 and T2 interleaves. This is not a strict schedule because T1 has an R(A), before T1 commits/aborts, W(A) in T2 is executed.

S3. The given schedule is a schedule, serial schedule and a strict schedule. It reads and writes are in the same order as in their original programs. T1 and T2 are executed one by one, without interleaving. T1 has a R(A), before T1 commits/aborts, W(A) in T2 is executed.

S4. The given schedule is not a schedule, not a serial schedule, and not a strict schedule. R(C) in P1 is missing in T1.

3. (60 points) Consider the following two schedules implemented by strict 2PL.

S1		S2	
T1	T2	T1	T2
S(A)		S(A)	
R(A)		R(A)	
X(A)			S(A)
W(A)			R(A)
commit		X(A)	commit
	S(A)		
	R(A)	W(A)	
	commit	commit	

Recall strict 2PL is implemented with a lock table to keep 1) which data object is locked; 2) which transaction owns which lock; 3) which transaction is waiting for a lock on this object.

Lock table	Data	Lock	Owner	Waiting

For each of the two schedules, explain how the status of the lock table changes as the actions are executed. The table is assumed to be empty initially.

Please use PowerPoints or other tools to type your answer. Don't handwrite. Submit your work through your Canvas account.

S1	
T1	T2
S(A) R(A) X(A) W(A) commit	S(A) R(A) commit

1. T1.S(A) needs a share lock on A. Since there isn't a lock on A, the application is approved.

Data	Lock	Owner	Waiting
A	S	T1	

2. T1.R(A) reads A, so the lock table remains the same

Data	Lock	Owner	Waiting
A	S	T1	

3. T1.X(A) needs an exclusive lock on A. Since A has a shared lock but this lock is owned by T1, the application is approved. The lock is upgraded to an exclusive lock.

Data	Lock	Owner	Waiting
A	X	T1	

4. T1.W(A) writes A. So, the lock table remains the same.

Data	Lock	Owner	Waiting
A	X	T1	

5. T1.commit. T1 finishes, so all locks it owns are now released.

Data	Lock	Owner	Waiting

6. T2.S(A) needs a shared lock on A. Because no lock on A exists, the application is approved.

Data	Lock	Owner	Waiting
A	S	T2	

7. T2.R(A) reads A. No change on the table.

Data	Lock	Owner	Waiting
A	S	T2	

8. T2.commit. T2 finishes, so all locks it owns are released

Data	Lock	Owner	Waiting

S2	
T1	T2
S(A) R(A)	
X(A)	S(A) R(A)
	commit
W(A) commit	

1. T1.S(A) needs a share lock on A. Since there isn't a lock on A, the application is approved.

Data	Lock	Owner	Waiting
A	S	T1	

2. T1.R(A) reads A, so the lock table remains the same

Data	Lock	Owner	Waiting
A	S	T1	

3. T2.S(A) needs a shared lock on A. Since A has a shared lock, the application is approved.

Data	Lock	Owner	Waiting
A	S	T1, T2	

4. T2.R(A) reads A. So the lock table remains the same.

Data	Lock	Owner	Waiting
A	S	T1, T2	

5. T1.X(A) needs an exclusive lock on A. Since A has a shared lock from T2, T1 is suspended.

Data	Lock	Owner	Waiting
A	S	T1, T2	T1(X)

6. T2.commit. T2 finishes and all locks it owns are released. T1 is still waiting, so it resumes. It needs an exclusive lock, which is not approved.

Data	Lock	Owner	Waiting
A	X	T1	

7. T1.W(A) writes A. No change on the table.

Data	Lock	Owner	Waiting
A	X	T1	

8. T1.commit. T1 finishes, so all locks it owns are released

Data	Lock	Owner	Waiting